

Assignment 2

Due: 9.00pm Tuesday 6th March 2018

Rules

1. This is a group assignment. (There are approximately 3 people per group and by now you should know your assigned group.)
 2. While **R** is the default package / programming language for this course you are free to use **R** or **Python** for the programming components of this assignment.
 3. Within each group **I strongly encourage each person to attempt each question by his / herself first** before discussing it with other members of the group.
 4. Students should **not** consult students in other groups when working on their assignments.
 5. Late assignments will **not** be accepted and all assignments must be submitted through the Hub with one assignment submission per group. Your submission should include a PDF report with your answers to each question as well as any relevant code. Make sure your PDF clearly identifies each member of the group by CID and name.
-

Question 1. Large-Scale Ridge Regression (30 points)

The ridge regression problem is given by

$$\min_{\boldsymbol{\beta}} \left\{ \|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|_2^2 + \lambda \|\boldsymbol{\beta}\|_2^2 \right\},$$

where $\boldsymbol{\beta} \in \mathbb{R}^{d+1}$,

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix} \quad \text{and} \quad \mathbf{X} = \begin{bmatrix} 1 & x_{11} & x_{12} & \dots & x_{1d} \\ 1 & x_{21} & x_{22} & \dots & x_{2d} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{N1} & x_{N2} & \dots & x_{Nd} \end{bmatrix} = \begin{bmatrix} 1 & \mathbf{x}_1^\top \\ 1 & \mathbf{x}_2^\top \\ \vdots & \vdots \\ 1 & \mathbf{x}_N^\top \end{bmatrix}.$$

The optimal solution of this problem is given by

$$\boldsymbol{\beta}^* = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^\top \mathbf{y}.$$

(In reality we don't shrink the constant coefficient β_0 in ridge regression (or Lasso) but to simplify notation here we assume that every coefficient is being shrunk.)

- (i) Suppose the dimension d of the data points is not large, but the number of data points N is *extremely* large. In fact N is so large that the matrix $\mathbf{X} \in \mathbb{R}^{N \times (d+1)}$ does not even fit in

memory. All of the data, however, is available in some database that can be queried for each data point. Explain how you could still compute the ridge estimate β^* . (10 points)

Hint: $\mathbf{X}^\top \mathbf{X}$ cannot be computed directly since \mathbf{X} is too large to store in memory. But perhaps $\mathbf{X}^\top \mathbf{X}$ can be written as a sum of the form $\sum_{i=1}^N \mathbf{z}_i^\top \mathbf{z}_i$ where \mathbf{z}_i is a $1 \times (d+1)$ vector ...?

- (ii) Suppose the data for the ridge regression problem becomes available sequentially, i.e. the k^{th} data point \mathbf{x}_k arrives at time t_k . At time t_k we want to be able to compute the optimal ridge estimate β_k^* using all the previous data $\{\mathbf{x}_1, \dots, \mathbf{x}_k\}$. In the usual method for ridge regression, we will have to store all the previous data points in a database, create the data matrix, and compute the estimate. Thus, the memory requirements will grow over time. Explain how you could still compute β_k^* for all $k \geq 1$, while keeping only $\mathcal{O}(d^2)$ numbers in the database. (10 points)
- (iii) Inverting a $(d+1) \times (d+1)$ matrix (like $(\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I})$) takes $\mathcal{O}(d^3)$ time so computing β_n^* (given a new observation \mathbf{x}_n) in part (b) is computationally expensive. Can you find a way of finding β_n^* that takes only $\mathcal{O}(d^2)$ time? (10 points)

Hint: Use the Sherman-Morrison-Woodbury identity $(\mathbf{A} + \mathbf{u}\mathbf{v}^\top)^{-1} = \mathbf{A}^{-1} - \frac{\mathbf{A}^{-1}\mathbf{u}\mathbf{v}^\top\mathbf{A}^{-1}}{1 + \mathbf{v}^\top\mathbf{A}^{-1}\mathbf{u}}$.

Question 2. Estimating the Scofflaw Rate in Chicago¹ (30 points)

A city that has parking meters on many of its streets wanted to estimate how much money it was losing because people were not paying the meters. Ideally, it would select a random sample of cars parked at meters at random times, and determine for each whether the meter read violation at the moment the car was approached. Then it would estimate p , the proportion of parked cars in violation of the meter, by \hat{p} , the fraction of such cars within the random sample.

However, such a sampling scheme was not practical. Even if the city knew exactly which meters had cars parked at them, a random selection among those cars might first pick a vehicle nine miles North of the city center, and then pick a second vehicle three miles South. Moving among the vehicles thus sampled would be hopelessly inefficient and expensive. Another scheme was necessary.

What the city decided to do instead was to pick metered blocks at random (e.g., 33rd Street between 3rd and 4th Avenues). Then, at a random time during a given interval (e.g., noon-1 PM), a surveyor would arrive at each such block and count both the number of cars parked there and the number of cars in violation of the meters. If a total of n blocks were sampled, and z_i was the number of cars parked on block i and x_i the number of “violation” cars there, then the overall “scofflaw rate” p would be estimated by \tilde{p} that follows:

$$\tilde{p} = \frac{\sum_{i=1}^n x_i}{\sum_{i=1}^n z_i}. \quad (1)$$

¹This question (and the next one) is taken from an excellent new textbook *Applied Statistics: Models and Intuition* (2015, Dynamic Ideas) by Arnold I Barnett. It provides an entertaining yet lucid introduction to statistics with real-world applications. If you’re interested in brushing up on your (classical) statistics and want to gain some real insight into how to think about real-world problems then this is a perfect book for you.

- (i) In words, what approximation is being made in (1)? (5 points)
- (ii) Why does that approximation seem a reasonable one? (5 points)
- (iii) In working out the margin of error, would it be reasonable to apply the usual formulas as if \tilde{p} were estimated by the fraction of violations within a purely random sample of $\sum_{i=1}^n z_i$ different cars? (In this case the “usual formulas” would say that an approximate 95% confidence interval (CI) for p is given by

$$\tilde{p} \pm 1.96 \times \sqrt{\frac{\tilde{p}(1 - \tilde{p})}{\sum_{i=1}^n z_i}} \quad (2)$$

for example.) Justify your answer. (10 points)

- (iv) Is \tilde{p} an unbiased estimator for p ? Justify your answer by either proving that it is unbiased or come up with a counter-example which shows that \tilde{p} is biased. (*Hint*: Imagine a city that has three blocks with parking meters, two of which will be sampled at random. Can you find a set of (x_i, z_i) ’s for which $\mathbb{E}[\tilde{p}] \neq p$? The example can be simple.) (10 points)

Question 3. Bootstrapping to Estimating the Scofflaw Rate in Chicago (35 points)

Continuing on from the previous question, suppose that the city samples 12 metered blocks at random, and reaches the following data of the form (x_i, z_i) :

(4,12)	(3,8)	(3,9)	(3,16)	(2,7)	(4,15)
(4,10)	(3,15)	(2,6)	(1,12)	(3,8)	(2,14)

- (i) What is the city’s estimate \tilde{p} of the citywide “scofflaw” rate p ? (5 points)
- (ii) Why is this a situation where bootstrapping seems desirable? (5 points)
- (iii) In a bootstrap analysis, what would be the approximate distribution of (X, Z) across metered blocks? (5 points)
- (iv) Perform a bootstrap analysis with 1,000 bootstrap samples, and assess whether the estimation procedure suffers an appreciable bias. (10 points)
- (v) Construct a 95% confidence interval for p based on percentiles in the bootstrap distribution of \tilde{p} ? (5 points)
- (vi) Create a histogram of the bootstrap statistic under resampling, and in the context of this histogram discuss why adjusting the bootstrap estimate of the “raw” confidence interval is not important here. (The raw confidence interval is the confidence interval $[q_l, q_u]$ where q_l and q_u are the $\alpha/2$ lower- and upper-sample quantiles, respectively, of the bootstrap samples of the test statistic.) (5 points)

Question 4. Comparing Classification Algorithms on the Tahoe Data-Set

(40 points)

Use logistic regression, classification trees, bagging and random forests to predict the 30-day readmission rate on the **Tahoe Healthcare Systems** data-set. In the case of logistic regression you should use a Lasso penalty term $\lambda \|\beta\|_1$ with λ then chosen via cross-validation. (This can be done using the `glmnet` package with the options $\alpha = 1$ (for Lasso) and `family = "binomial"` (for logistic regression). For those of you who are interested, the model is fit by via maximizing the log-likelihood of the data + the regularization term, $\lambda \|\beta\|_1$.) For each of the 4 methods using the same data split with 80% of the data assigned to the training set and 20% to the test set. Use 5-fold cross-validation on the training set to choose appropriate parameter values, e.g. λ for logistic regression with Lasso penalty, m for random forests etc.

Report your results. Is there a clear winner in terms of performance on the test set?
